

Python Test

Project Description:

Build a RAG-based Question-Answering System

Objective:

Create a Retrieval-Augmented Generation (RAG) API using an LLM to answer questions over a set of documents. The system should be scalable, API-driven, and containerised.

Key requirements:

- **LLM Integration:**
 - Use an LLM API (e.g., OpenAI GPT, Gemini) or an open-source model.
 - Connect it with a retrieval pipeline (e.g., FAISS, Weaviate, Chorma).
- **Backend API:**
 - Build APIs using FastAPI or Flask.
 - Expose endpoints for:
 - Document Upload & Indexing
 - Query Answering
 - Fetching Logs
- **Logging:**
 - Store queries and responses in a NoSQL database (e.g., MongoDB).
 - Log query text, context, response, timestamp, and duration.
 - Provide an endpoint to fetch logs.
- **Deployment:**
 - Dockerize the project for easy deployment.
 - Include a Dockerfile and clear setup instructions.
- **Optional (Bonus Points):**
 - Implement a multi-agent workflow (e.g., retrieval agent + refinement agent).
 - Add basic monitoring and health-check APIs.

Deliverables:

- **Code Repository:**
 - Source code with APIs, logging, and documentation.
 - Dockerfile for deployment.
 - README with setup instructions.
- **NoSQL Logs:**
 - Sample log entries stored in the database.
 - API endpoint to retrieve logs.
- **Optional:**
 - Multi-agent workflow example.
 - Monitoring or health-check endpoints.

Evaluation Criteria:

API Functionality
Query Accuracy
NoSQL Logging Implementation
Code Quality & Documentation
Docker Deployment
Optional Features (Bonus)